


CS 7180: Agile/Scrum + Pair Workflow



Northeastern
University

John Alexis Guerra Gomez

jguerra at northeastern.edu

Class: johnguerra.co/classes/aiCoding_spring_2026

Slides: johnguerra.co/lectures/ai_assisted_coding

What We'll Cover Today

- 1. Where We Are -- Week 7 checkpoint
- 2. Agile/Scrum Crash Course
- 3. GitHub as Your Scrumboard
- 4. From PRD to Sprint Backlog
- 5. Branches, PRs, and Code Review
- 6. Pair Workflow with AI
- 7. Code Review Between Partners

Where We Are

| *Week 7 -- Process for working in pairs*

Recap: Week 6 Foundations

Last week you learned:

- How IDE AI tools work (context collection, indexing, embeddings)
- Code suggestions, inline edit (Cmd+K), chat panel
- Modes: Ask / Write / Agent / Plan
- Rules files and @ context references
- Tool comparison: Antigravity vs Copilot vs Cursor
- **P2 pair formation** -- you should have a partner and a Canvas group
You have the tools and a partner. Now you need the *process* to work effectively together.

This Week: Process for Pairs

Two themes today:

- 1. **Agile/Scrum** -- How professional teams organize work (you'll need this for P2 and P3)
- 2. **Pair workflow** -- How to split work, avoid conflicts, and review each other's code

Agile gives you the *structure*, your partner gives you *accountability*, and AI gives you *speed*.

Agile/Scrum Crash Course

The minimum you need to run effective sprints

Why Agile for AI Projects?

AI changes the speed of coding, not the need for process.

- You can generate code faster than ever -- but *what* should you build?
- Without process, AI speed leads to building the wrong thing faster
- Agile gives you **short feedback loops** to course-correct
- Perfect fit: sprints align with project milestones in this course

P2 requires 2 documented sprints. P3 requires 4.

Scrum Roles

Role	Responsibility	In This Course
Product Owner	Decides <i>what</i> to build, prioritizes backlog	One partner takes lead (rotate per sprint)
Scrum Master	Facilitates process, removes blockers	The other partner (rotate per sprint)
Development Team	Builds the product	Both partners + your AI tools

In P2, you split roles between partners. In P3 teams, you'll have more flexibility.

The Sprint Cycle

```
Sprint Planning (start of sprint)
↓
Daily Work (build, test, commit)
↓
Sprint Review (demo what you built)
↓
Sprint Retrospective (what to improve)
↓
Next Sprint Planning...
```

Sprint length in this course: 1-2 weeks per sprint.

The key insight: you commit to a *small, achievable* set of work each sprint.

Scrum Ceremonies

Ceremony	When	Duration	Purpose
Sprint Planning	Start of sprint	30-60 min	Pick issues for this sprint
Daily Standup	Every day	5-15 min	What did I do? What will I do? Blockers?
Sprint Review	End of sprint	30 min	Demo working software
Retrospective	After review	15-30 min	What went well? What to improve?

For P2: Document your planning and retro in your README or project wiki.
Partner standups count -- at least 3 per sprint.

GitHub as Your Scrumboard

Issues, Projects, and Labels -- everything you need

GitHub Issues = Backlog Items

Each piece of work becomes a GitHub Issue:

- **Title:** Short, action-oriented ("Add user login page")
- **Description:** Acceptance criteria, design notes, links
- **Labels:** feature, bug, chore, sprint-1, sprint-2
- **Assignee:** Who's working on it
- **Milestone:** Which sprint or release

Issues are your single source of truth for what needs to be built.

GitHub Projects = Sprint Board

GitHub Projects (Board view) gives you a Kanban board:

Backlog	Sprint Todo	In Progress	In Review	Done
Future work	This sprint's work	Actively coding	PR open	Merged & deployed

Setup:

1. Create a Project in your repo (Board layout)
2. Add columns: Backlog, Todo, In Progress, Review, Done
3. Drag issues between columns as work progresses
4. Filter by sprint milestone to see current sprint only

Labels & Milestones for Sprint Tracking

Labels categorize work:

- feature / bug / chore / docs
- priority: high / priority: low
- sprint-1 / sprint-2

Milestones group issues into time-boxed sprints:

- Create a milestone for each sprint with a due date
- Assign issues to milestones during sprint planning
- Track progress via the milestone's completion percentage

For P2: Create at least 2 milestones (Sprint 1, Sprint 2). Each should have 5+ issues.

From PRD to Sprint Backlog

Connecting what you learned in Weeks 3-4 to how you build in Weeks 7+

Revisiting the PRD

Remember your PRD from Weeks 3-4? It contains:

- Problem statement and target users
- User stories ("As a ___, I want ___ so that ___")
- Acceptance criteria for each story
- Technical architecture decisions

Now it's time to turn that PRD into actionable work items.

Your PRD is the *what*. Your sprint backlog is the *when* and *how*.

Breaking PRD into GitHub Issues

Each user story becomes one or more GitHub Issues:

PRD User Story:

"As a user, I want to log in with email
so that my data is saved across sessions"

GitHub Issues:

- #1 Set up authentication library (chore)
- #2 Create login page UI (feature)
- #3 Implement email/password auth endpoint (feature)
- #4 Add session persistence (feature)
- #5 Write login flow tests (chore)

Rule of thumb: Each issue should be completable in 1-4 hours. If it's bigger, break it down further.

Sprint Planning: Picking Issues

In sprint planning, you and your partner:

1. Review the backlog together (all open issues)
2. Estimate effort (small / medium / large)
3. Pick issues that fit the sprint's capacity
4. **Assign each issue to one partner** -- clear ownership
5. Assign them to the sprint milestone

For Sprint 1 of P2, aim for:

- Core data model and API setup
- 1-2 key features (login, main CRUD operation)
- Basic test suite setup
- Shared rules file established

Branches, PRs, and Code Review

Professional workflow with AI assistance

Branch-per-Issue Workflow

Every issue gets its own branch:

```
# Create branch from issue number  
git checkout -b feature/42-add-login-page  
  
# Work on the feature...  
git add .  
git commit -m "Add login page UI (#42)"  
  
# Push and create PR  
git push -u origin feature/42-add-login-page
```

Naming convention: type/issue#-short-description

- feature/42-add-login-page
- fix/57-null-avatar-crash
- chore/63-update-dependencies

PRs and AI-Assisted Review

Pull Request best practices:

- **Link to issue:** "Closes #42" in the PR description auto-closes the issue on merge
- **Small PRs:** Easier to review, fewer conflicts, faster to merge
- **PR template:** Add a template in `.github/pull_request_template.md`
- **AI-assisted review:** Use your IDE AI to review diffs before pushing

AI-assisted self-review before creating PR:

"Review this diff for bugs, security issues, and style violations.
Our project uses TypeScript, React, and follows the patterns in
`@.antigravityrules`"

Pair Workflow with AI

Scrum for two -- agile process meets AI-assisted development

The Pair + AI Workflow

Your pair runs scrum, AI assists both partners:

Activity	How It Works	AI's Role
Sprint Planning	Meet together, pick issues from backlog	Helps break down stories into tasks
Development	Each partner owns assigned issues	Real-time coding assistance per partner
Code Review	Every PR reviewed by your partner	Pre-review via AI, human catches intent
Standups	Async check-ins on progress & blockers	Summarize changes, draft standup notes

When to pair vs. split work:

- **Pair (synchronous):** Complex features, architecture decisions, debugging hard issues
- **Split (async):** Independent features with clear interfaces, tests, documentation

Your **design thinking** and **mom test** skills feed the backlog -- user feedback drives sprint priorities.

Splitting Work Without Conflicts

How to avoid stepping on each other's toes:

1. **Assign issues to one person** -- never both working on the same file
2. **Define interfaces first** -- agree on API shapes, component props, data models before splitting
3. **Use feature branches** -- one branch per issue, never commit directly to main
4. **Merge frequently** -- don't let branches diverge for days
5. **Communicate blockers** -- async standups keep both partners aware

Anti-pattern: Both partners editing the same file in parallel. This leads to merge conflicts and wasted time.

Async Standups for Pairs

Not everyone can meet daily. Use async standups:

Post in your shared channel (Slack DM, GitHub Discussion, or project wiki):



Date: [today]

Yesterday: Completed #42 login UI, started #43 auth endpoint



Today: Finish auth endpoint, write tests for login flow



Blockers: Waiting on partner's DB schema PR (#40) to merge

For P2: At least 3 standups per sprint from each partner. These can be async messages.

Code Review Between Partners

| *Every PR gets a human review*

Partner Code Review Workflow

Every PR reviewed by your partner before merge:

1. Developer creates PR (links to issue)
2. Partner reviews the code
 - Read the diff
 - Check against acceptance criteria
 - Run locally if needed
 - Leave comments (questions, suggestions, approvals)
3. Developer addresses feedback
4. Partner approves → Merge

Minimum 5 PR reviews per partner across the project (visible in GitHub).

What to Look For in Review

When reviewing your partner's PR:

- **Does it match the issue?** -- Check acceptance criteria
- **Code quality** -- Naming, structure, readability
- **Tests included?** -- New feature should have tests
- **Rules file followed?** -- Does it match your shared conventions?
- **No dead code** -- Remove console.logs, commented-out code

Use AI to assist your review but don't skip reading the code yourself. AI catches patterns; humans catch intent.

What to Remember

- 1. **Agile isn't bureaucracy** -- it's short feedback loops to build the right thing faster.
- 2. **GitHub Issues + Projects = your scrumboard** -- no extra tools needed.
- 3. **PRD maps to issues** -- every user story becomes actionable GitHub Issues.
- 4. **Branch-per-issue** keeps your work organized and reviewable.
- 5. **Every PR gets reviewed by your partner** -- this is how professionals work.
- 6. **Pair workflow + AI** = scrum process, design thinking, and AI working together.

Looking Ahead

Next Week: Advanced IDE AI Features

Week 8 -- Power features for your P2 sprint

- **Agent memory** -- Persistent project knowledge via `.antigravityrules` / `CLAUDE.md`
 - **MCP servers** -- Connecting your AI to external tools and data
 - **Browser mode & mockup-to-code** -- AI that sees your running app
 - **Debugging with AI** -- Error analysis, stack trace reading, rubber duck debugging
 - **Shared rules files for pairs** -- Evolving conventions as a team
- HW3 (Context Engineering) is due Week 8.** Create your P2 rules file and Scrum board. Start now.

Resources

Required Reading

Resource	URL
The Scrum Guide (official)	<u>scrumguides.org</u>
GitHub Projects Documentation	<u>docs.github.com/issues/planning-and-tracking-with-projects</u>
GitHub Issues Documentation	<u>docs.github.com/issues</u>

Recommended Reading

Resource	URL
Scrum by Jeff Sutherland (course textbook)	Required book for the course
Atlassian Agile Coach	<u>atlassian.com/agile</u>
GitHub Flow Guide	<u>docs.github.com/get-started/using-github/github-flow</u>

Speaker notes